Compiled by: Adarsha Paudyal

AWS HAND-ON-GUIDE

Welcome to this comprehensive guide on AWS (Amazon Web Services). In this document, you'll find 25 solved AWS sample questions to help you better understand the AWS.

1. How does cloud computing impact the overall cost structure for businesses compared to traditional IT infrastructure?

Cloud Computing and Its Impact on Cost Structure

- Cloud Computing Defined
- ▲ On-demand access to computing power, databases, storage, and applications.
- **3** Pay-as-you-go pricing—only pay for what you use.
- To Delivered via the **internet**, no physical hardware needed.

③ Cost Structure: Traditional IT vs. Cloud Computing

- Traditional IT (Infrastructure as Hardware)
- 🚧 High Upfront Costs Buying servers, storage, networking equipment.
- **Requires Space & Staff** Data centers need physical space, security, and maintenance.
- 🐌 Slow Scaling Long procurement cycles and guessing future capacity needs.
- ▲ Cloud Computing (Infrastructure as Software)
- Flexible & Scalable Instantly add or remove resources based on demand.
- **Faster Deployment** No waiting for hardware procurement.
- 🂸 Lower Costs No need for data center management, maintenance, or excess capacity.
- 🚀 Eliminates Heavy Lifting Focus on business innovation, not infrastructure.
- Cloud = Lower Costs, Faster Growth, More Agility! 🚀
 - 2. Discuss the Key Benefits of Adopting Cloud Computing with a Suitable Example.

Key Benefits of Adopting Cloud Computing (Easy-to-Remember Format)

1. Trade Capital Expense for Variable Expense

- Before: Huge upfront investment in data centers.
- After: Pay only for the resources you use, like a utility bill.

2. Massive Economies of Scale

- Cloud providers like AWS serve millions of customers.
- More usage = Lower costs for everyone.

3. Stop Guessing Capacity

Problem:

- Overestimate → Wasted money on unused servers.
- Underestimate → Performance issues.
 - Solution: Cloud scales up or down on demand.

4. Increase Speed and Agility

- Before: Weeks or months to get new servers.
- After: Deploy resources in minutes.

5. Stop Spending on Data Centers

No more costs for:

- ✓ Landscaping
- ✓ Hardware

Savings → Invest in Business Growth

6. Go Global in Minutes

- Deploy applications worldwide with just a few clicks.
- Reach customers across different regions instantly.

Example: E-commerce Website Deployment

- Buy multiple servers and set up a data center.
- Hire IT staff for maintenance and security.
- Guess traffic demand (risk of over/under-provisioning).
- Servers may **crash** due to overload.

▲ After Cloud (Cloud Computing)

Use AWS Elastic Beanstalk or Google Cloud App Engine.

- Auto-scales up/down based on demand.
- Minimal cost on normal days, scales up during high traffic.
- No crashes handles millions of visitors smoothly.

✓ Cloud = Faster, Scalable, Cost-Effective Deployment!

3. How do cloud service models such as SaaS, PaaS, and IaaS differ from each other?

Difference Between SaaS, PaaS, and IaaS

Software as a Service (SaaS) – Ready-to-Use Software

- ✓ What it is: Pre-built software, accessible via the internet.
- ✓ For Users No need to install or manage anything.
- ✓ Flexible? Yes, access from any device with an internet connection.
- ✓ Scalability? Highly scalable, as the provider manages everything.
- ✓ Maintenance? No maintenance required; the provider updates automatically.
- ✓ Examples: Gmail, Google Drive, Zoom, Netflix.

* Platform as a Service (PaaS) – For Developers

- ✓ What it is: A platform to develop, test, and deploy apps.
- ✓ **For Developers** Focus on coding; no need to manage hardware.
- ✓ Flexible? Yes, allows custom app development.
- ✓ Scalability? Scales based on application needs.
- ✓ **Maintenance?** Partially managed; developers handle code, but infrastructure is managed by the provider.
- ✓ Examples: AWS Elastic Beanstalk, Google App Engine, Heroku.

Infrastructure as a Service (IaaS) – Full Control Over Resources

- ✓ What it is: Cloud-based virtual servers, storage, and networking.
- ✓ For IT Admins Full control over infrastructure.
- ✓ **Flexible?** Yes, configure your own servers and network.
- ✓ Scalability? Highly scalable but requires manual configuration.
- ✓ Maintenance? Users are responsible for managing and securing the infrastructure.
- ✓ Examples: AWS EC2, Google Compute Engine, Microsoft Azure VM.

Feature	SaaS 🖥	PaaS 🔧	laaS 🌣
Who Uses It?	End-users (non-tech)	Developers	IT Admins & Developers
Control Level	Least control	Medium control	Full control

Feature	SaaS 🖥	PaaS 🔧	laaS 🌣
Flexibility	Limited (pre-built software)	Custom app development	Fully customizable
Maintenance	No maintenance	Some maintenance	Full responsibility
Examples	Gmail, Netflix	Heroku, App Engine	AWS EC2, Azure VM

- SaaS → Software → Just Use It (Netflix, Gmail).
- PaaS → Platform → Build & Deploy Apps (Heroku, App Engine, Vercel).
- IaaS → Infrastructure → Full Control Over Servers (AWS EC2, Azure VM).
- 4. Discuss various Cloud Computing Deployment Models with a Suitable Example.

Cloud Computing Deployment Models

1. Public Cloud

- **Description:** Services offered over the internet by third-party providers; resources are shared among multiple organizations.
- **Pros:** Cost-effective, scalable, and easy to deploy.
- Cons: Less control over security and customization.
- **Example:** A startup uses Amazon Web Services (AWS) to host its applications, benefiting from AWS's infrastructure without investing in physical servers.

2. Private Cloud

- **Description:** Exclusive cloud environment dedicated to a single organization, either on-premises or hosted by a third-party.
- **Pros:** Enhanced security, greater control, and customization.
- **Cons:** Higher costs and maintenance responsibilities.
- **Example:** A financial institution operates a private cloud to manage sensitive customer data, ensuring compliance with regulatory standards.

3. Hybrid Cloud

- **Description:** Combines public and private clouds, allowing data and applications to be shared between them.
- **Pros:** Flexibility, optimized workload distribution, and cost efficiency.
- **Cons:** Complex management and potential security challenges.
- **Example:** A retail company uses a private cloud for its customer databases while leveraging the public cloud for handling increased traffic during holiday sales.

4. Community Cloud

• **Description:** Shared cloud infrastructure among several organizations with common concerns, such as security or compliance.

- **Pros:** Cost-sharing, collaborative environment, and tailored services.
- Cons: Limited control and potential for conflicts between organizations.
- **Example:** Multiple healthcare providers collaborate on a community cloud to manage patient records, ensuring adherence to healthcare regulations.

Easy to Remember:

- Public Cloud: Shared resources, accessible to all.
- Private Cloud: Exclusive resources for one organization.
- Hybrid Cloud: A mix of public and private clouds.
- Community Cloud: Shared among organizations with common interests.
- 5. Discuss Six AWS Cloud Adoption Framework with a Suitable Example.

1. Business Perspective

- Align IT with business goals.
- Ensure IT investments deliver measurable results.
- Manage business risks effectively.
- Focus on financial and strategic goals.
- **Key Roles**: Business Managers, Finance Managers.

2. People Perspective

- Develop skills and provide training.
- Manage career growth and incentives.
- Support organizational change.
- Build an agile and adaptable workforce.
- **Key Roles**: HR Managers, Staffing Managers.

3. Governance Perspective

- Align IT strategy with business strategy.
- Implement risk management and compliance.
- Manage projects and business performance.
- Focus on portfolio and program management.
- **Key Roles**: CIO, Program Managers, Business Analysts.

4. Platform Perspective

- Manage compute, network, storage, and databases.
- Design and implement scalable solutions.

- Migrate and modernize infrastructure.
- Focus on application development and architecture.
- Key Roles: CTO, IT Managers, Solution Architects.

5. Security Perspective

- Protect data with strong security measures.
- Implement access controls and monitoring.
- Ensure compliance and security policies.
- Respond quickly to security incidents.
- Key Roles: CISO, IT Security Managers.

6. Operations Perspective

- Optimize service and application performance.
- Monitor resources and manage change.
- Ensure business continuity and disaster recovery.
- Automate operations to improve efficiency.
- **Key Roles**: IT Operations Managers, IT Support Managers.

6. Buddha Air is planning to move its entire Infrastructure from On-premises to Cloud based Infrastructure. Identify and compare the Total Cost of Ownership (TCO) for on-premises and cloud based infrastructure. The Total Cost of Ownership should include Server Costs, Storage Cost, Network Cost and IT Labor Cost.

Here is the **detailed 3-year Total Cost of Ownership (TCO) breakdown** for **On-Premises vs. AWS Cloud** for Buddha Air,

1. On-Premises TCO Breakdown (3 Years)

© On-premises infrastructure requires significant upfront investment in hardware, network, and IT labor, along with ongoing maintenance and operational costs.

1. Server Cost

- Initial Costs: \$75,000 (Physical servers, installation, power setup)
- Maintenance Costs (Over 3 Years): \$16,922 (Cooling, power, hardware replacement at ~7.5% annually)
- Total (3 Years): \$91,922

2. Storage Cost

- Initial Costs: \$50,000 (Storage Area Network (SAN), NAS, backup solutions)
- Maintenance Costs (Over 3 Years): \$17,840 (Storage expansion, data redundancy, security)
- Total (3 Years): \$67,840

3. Network Cost

- **Initial Costs**: \$5,000 (Firewalls, routers, switches, load balancers)
- Bandwidth Costs (3 Years): \$3,600 (~\$100/month for a dedicated connection)
- Maintenance Costs (Over 3 Years): \$1,060 (Upgrades, replacements)
- Total (3 Years): \$7,660

4. IT Labor Cost (Including Training & Maintenance)

- **Training Cost**: \$2,000 (Cloud migration training, cybersecurity, software maintenance)
- Total (3 Years): \$2,000

On-Premises Total Cost of Ownership (TCO) Over 3 Years

§ \$91,922 (Server) + \$67,840 (Storage) + \$7,660 (Network) + \$2,000 (IT Labor) = \$167,422

2. AWS TCO Breakdown (3 Years)

AWS operates on a pay-as-you-go model, significantly reducing capital expenses and eliminating maintenance costs.

1. Server Cost (AWS EC2 Compute)

- Monthly Cost: \$70 (t3.medium instance)
- Total (3 Years): \$2,547

2. Storage Cost (AWS S3 & Backup)

- Monthly Cost: ~\$140 (3TB on Amazon S3 + backup storage)
- Total (3 Years): \$4,963

3. Network Cost (Bandwidth, Load Balancer, VPN, Outbound Data)

- AWS covers most networking costs
- Total (3 Years): Included

4. IT Labor Cost (Managed Services Reduce Labor Needs)

- · AWS eliminates most maintenance and training needs
- Total (3 Years): Included

AWS Total Cost of Ownership (TCO) Over 3 Years

💰 \$2,547 (Server) + \$4,963 (Storage) + \$0 (Network) + \$0 (IT Labor) = \$7,509

3-Year Total Cost of Ownership (TCO) Comparison

Category On-Premises AWS Cloud

Category	On-Premises	AWS Cloud
Servers (Hardware, Compute)	\$91,922	\$2,547
Storage (SAN/NAS, Backups, Cloud Storage)	\$67,840	\$4,963
Network (Bandwidth, Load Balancing, VPN)	\$7,660	Included
IT Labor (Maintenance, Admin, Training)	\$2,000	Included
Total 3-Year TCO	\$167,422	\$7,509
Total Savings (96%)		\$159,913

Key Takeaways:

- By migrating to AWS Cloud, Buddha Air can save up to 96% of the total cost.
- The **3-year total savings would be \$159,913** compared to on-premises infrastructure.
- AWS eliminates costs for network infrastructure, IT labor, and maintenance, reducing the operational burden.
- AWS's pay-as-you-go model ensures you only pay for what you use, avoiding large upfront capital expenditures.
- **Cloud-Based Infrastructure** saves on IT labor costs and provides better **scalability & flexibility**, making it **more cost-effective** in the long run.
- On-Premises is more expensive due to high upfront costs, maintenance, and IT staffing but provides complete control over infrastructure.
- 7. Discuss the importance of Identity and Access Management in Cloud Computing with a Suitable Example.

Importance of Identity and Access Management (IAM) in Cloud Computing

IAM is a **security framework** that helps manage and control access to cloud resources. It ensures that **only authorized users** can access specific services, reducing security risks.

How IAM Helps in Cloud Computing?

- Manages cloud resources like servers, databases, and applications.
- Creates and controls AWS users and groups for access management.
- Assigns permissions based on roles to follow the least privilege principle.
- Enhances security with features like Multi-Factor Authentication (MFA) and access logs.

Four Key Components of AWS IAM

1. IAM Users 👤

- Represents individual users (e.g., developers, admins) who need access to AWS.
- Each user has unique **login credentials** (username + password or access keys).
- Users can have **specific permissions** based on policies.

2. IAM Groups 👥

- A collection of IAM users with common permissions.
- Easier to manage access by assigning **permissions to the group** instead of individual users.

Example: A company creates a "Developers Group" with access to EC2 and RDS. Any new developer added to the group gets the same access.

3. IAM Roles 🔄

- **Used for temporary access** by AWS services or external users.
- Unlike users, roles **don't have login credentials** (password or access keys).
- Assigned to AWS resources like **EC2 instances**, **Lambda**, **or third-party apps**.

✓ Example: An EC2 instance needs access to S3 storage. Instead of giving it a username/password, an IAM role is assigned with S3 read access.

4. IAM Policies 📜

- **JSON-based rules** that define **permissions** for users, groups, or roles.
- Specifies what actions are allowed or denied on AWS services.
- Policies are **attached to users, groups, or roles** to control access.
- ✓ Example: A policy can allow a user to read from S3 but deny write access.

IAM is a crucial part of **cloud security**, ensuring controlled and secure access to AWS services. By using **users, groups, roles, and policies**, businesses can effectively **manage permissions** and protect their cloud environment.

8. Discuss the Shared Security Model in Cloud Computing with a Suitable Example.

9. You have a small business with a website that is hosted on an Amazon Flastic

Compute Cloud (Amazon EC2) instance. You have customer data that is stored on

- **a.** Your web server and database server must be in separate subnets.
- **b.** The first address of your network must be 10.10.1.0. Each subnet must have 254 total IPv4 addresses.
- **c.** Your customers must always be able to access your web server.
- **d.** Your database server must be able to access the internet to make patch updates.
- **e.** Your architecture must be highly available and use at least one custom firewall layer.
- **f.** Design the architecture with a brief explanation of each service used.

To design an Amazon Virtual Private Cloud (VPC) architecture that meets your specified requirements, follow these steps:

1. VPC Creation:

- **CIDR Block**: Create a VPC with the CIDR block 10.10.1.0/24, which provides 256 IP addresses.
- Subnets:
 - **Public Subnet**: Create a public subnet with the CIDR block 10.10.1.0/25 (128 IP addresses).
 - Private Subnet: Create a private subnet with the CIDR block 10.10.1.128/25 (128 IP addresses).

2. Subnet Allocation:

- Web Server: Launch your web server in the public subnet.
- **Database Server**: Launch your database server in the private subnet.

3. Internet Access:

- **Web Server**: Ensure the web server in the public subnet has a public IP address and is associated with an Internet Gateway (IGW) to allow direct internet access.
- **Database Server**: Configure a Network Address Translation (NAT) Gateway in the public subnet to enable the database server in the private subnet to access the internet for patch updates.

4. Route Tables:

- Public Route Table:
 - **Association**: Associate this route table with the public subnet.
 - Routes:
 - Add a route for 0.0.0.0/0 targeting the Internet Gateway to allow internetbound traffic.
- Private Route Table:
 - **Association**: Associate this route table with the private subnet.
 - Routes:
 - Add a route for 0.0.0.0/0 targeting the NAT Gateway to allow internet-bound traffic from the private subnet.

5. Network Access Control Lists (NACLs):

- Public NACL:
 - Inbound Rules:

• Allow inbound HTTP (port 80) and HTTPS (port 443) traffic from the internet.

Outbound Rules:

• Allow all outbound traffic to the internet.

Private NACL:

Inbound Rules:

 Allow inbound traffic from the public subnet on the database port (e.g., port 3306 for MySQL).

Outbound Rules:

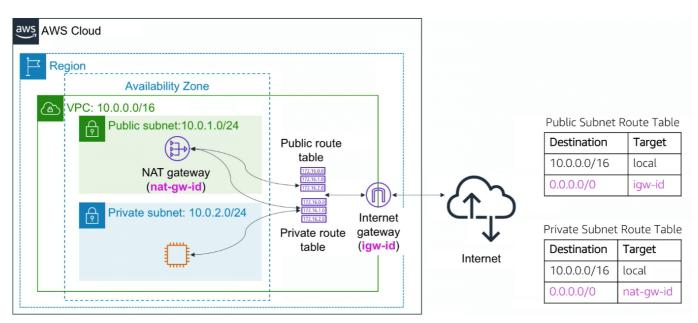
• Allow outbound traffic to the internet via the NAT Gateway.

6. High Availability:

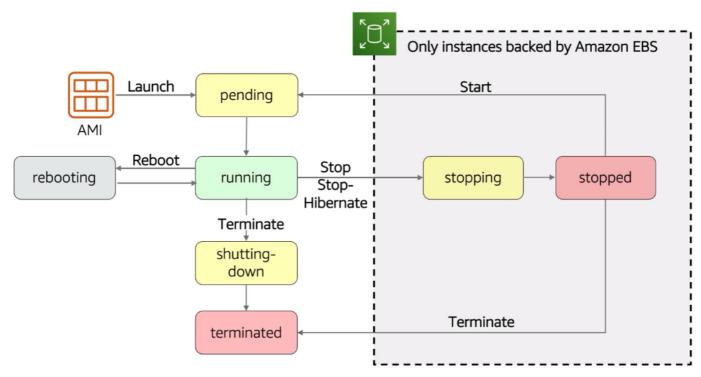
- **Multiple Availability Zones (AZs)**: Deploy the VPC across at least two AZs to ensure high availability.
- **Redundant NAT Gateways**: Deploy NAT Gateways in each AZ to provide fault tolerance for internet access from the private subnet.

7. Security Measures:

- Security Groups:
 - Web Server Security Group:
 - Allow inbound HTTP (port 80) and HTTPS (port 443) traffic from the internet.
 - Database Server Security Group:
 - Allow inbound traffic only from the web server's security group on the database port (e.g., port 3306 for MySQL).



10. Discuss the Amazon EC2 Instance Lifecycle with a suitable diagram.



Here's a more concise and simplified explanation of the **Amazon EC2 Instance Lifecycle** based on the diagram you shared:

1. Pending

- State: EC2 instance is being provisioned.
- What Happens: Resources are being allocated, like CPU and memory.
- **Billing**: Begins as soon as the instance starts.
- **Transition**: Moves to **Running** once provisioning is complete.

2. Running

- State: The instance is fully operational.
- What Happens: You can access the instance, run applications, and manage workloads.
- Possible Actions:
 - Reboot: Restarts without affecting the instance.
 - Stop: Shuts down while retaining EBS data.
 - Terminate: Permanently deletes the instance.
 - Hibernate: Saves RAM state and shuts down.
- Transition:
 - Can move to **Stopping**, **Rebooting**, or **Terminating**.

3. Stopping

- **State**: Instance is being shut down.
- What Happens: The instance is stopped, but data on EBS volumes is preserved.
- **Transition**: Moves to **Stopped** once shut down.

4. Stopped

- State: Instance is offline.
- What Happens: You can restart, modify, or terminate the instance.
- **Billing**: No compute charges, but EBS storage charges apply.
- Transition: Moves to Pending if restarting, or Terminated if deleted.

5. Rebooting

- State: Instance is being rebooted.
- What Happens: Only the OS restarts, the instance remains active.
- Transition: Back to Running after reboot.

6. Shutting-down

- State: Instance is being permanently terminated.
- What Happens: EBS volumes are deleted unless specified otherwise.
- Transition: Moves to Terminated once the shutdown process is complete.

7. Terminated

- State: Instance is permanently deleted.
- What Happens: All resources are released.
- **Billing**: Stops as soon as the instance is terminated.

8. Hibernated (Special Case)

- State: Instance is stopped, but RAM content is saved to EBS.
- What Happens: The instance can resume from its saved state later.
- **Transition**: Moves to **Stopped** after hibernation and to **Pending** when restarting.

11. Discuss the Four Pillars of Cost Optimization for Compute Services with a suitable example.

Apologies for the earlier confusion. Let's revisit the **Four Pillars of Cost Optimization for Compute Services** with clear explanations and practical examples:

- 1. **Right Size**: Choose the appropriate size for your compute resources to match your workload needs.
 - **Avoid Over-Provisioning**: Don't allocate more resources than necessary; this leads to unnecessary costs.
 - **Prevent Under-Utilization**: Ensure resources are used efficiently to avoid paying for unused capacity.

• **Regularly Assess Needs**: Continuously evaluate your workload requirements to adjust resource sizes accordingly. Example: If your application requires minimal CPU and memory, opting for a smaller instance type can reduce costs.

- 2. **Increase Elasticity**: Utilize the cloud's ability to scale resources up or down based on demand.
 - Auto-Scaling: Automatically adjust the number of running instances to match traffic levels.
 - **Dynamic Resource Allocation**: Allocate resources in real-time to handle varying workloads efficiently.
 - **Cost Efficiency**: Scale down during low demand periods to save costs. Example: Set up autoscaling groups to add more instances during high traffic periods and remove them during low traffic.
- 3. **Optimal Pricing Model**: Select the most cost-effective pricing options for your workloads.
 - **Reserved Instances**: Commit to using specific instance types over a term to receive discounted rates.
 - **Spot Instances**: Bid for unused capacity at lower prices for flexible, interruptible workloads.
 - Savings Plans: Purchase compute usage plans to save on costs over a set period. Example:
 Purchase Reserved Instances for predictable workloads to save up to 75% compared to On-Demand pricing.
- 4. **Optimize Storage Choices**: Select storage solutions that balance performance needs with cost considerations.
 - **Appropriate Storage Classes**: Use storage classes that align with data access frequency and performance requirements.
 - **Data Lifecycle Policies**: Automatically transition data to more cost-effective storage tiers as it ages.
 - Regular Data Audits: Periodically review stored data to remove unnecessary or outdated information. Example: Store infrequently accessed data in Amazon S3 Glacier to reduce storage costs.
- 12. Differentiate Containers with Virtual Machines with a suitable architecture.

Difference Between Containers and Virtual Machines (VMs) in AWS Cloud Computing

Feature	Containers	Virtual Machines (VMs)	
Definition	A lightweight, isolated environment that runs applications using a shared OS kernel.	A fully isolated operating system instance running on a virtualized server.	
Architecture	Uses a container runtime (e.g., Docker, Kubernetes) on a shared OS.	Uses a hypervisor (e.g., VMware, KVM) to run multiple OS instances.	

Feature	Containers	Virtual Machines (VMs)	
Isolation	Process-level isolation; all containers share the same OS kernel.	Full OS-level isolation; each VM has its own OS and resources.	
Resource Usage	Lightweight; consumes fewer resources.	Heavy; requires more CPU, RAM, and storage due to OS overhead.	
Startup Time	Fast (seconds) since it doesn't need to boot an OS.	Slower (minutes) as it loads a full OS for each VM.	
Operating System	Shares the host OS kernel (all containers must use compatible OS).	Can run different OS types (e.g., Linux and Windows on the same host).	
Portability	Highly portable across environments (local, cloud, hybrid).	Less portable due to OS dependencies.	
Management	Managed using container orchestration tools like Kubernetes, ECS, or EKS.	Managed using VM orchestration tools like AWS EC2, VMware, or OpenStack.	
Security	Less isolated than VMs, but security can be improved with proper policies.	More secure due to full OS-level isolation.	
Persistence	Containers are ephemeral (destroyed when stopped unless configured with persistent storage).	VMs have persistent storage by default.	
Use Case	Best for microservices, cloud-native apps, CI/CD, and serverless workloads.	Suitable for legacy applications, multi-OS environments, and workloads requiring strict isolation .	
AWS Services	Amazon ECS (Elastic Container Service), Amazon EKS (Elastic Kubernetes Service), AWS Fargate.	Amazon EC2 (Elastic Compute Cloud), VMware Cloud on AWS.	

Difference Between Containers and Virtual Machines (VMs) – Easy to Remember

ls) 🖥
virtualized server.
s its own OS).
. (1)
ypes (Windows, Linux,
cure. 🛭
Лware, Нурег-V.

Wishual Machines (WMs)

Feature	Containers T #		Virtual Machi	ines (VMs)	
Best For	Microservices, cloud apps, fast scaling.		Heavy worklo	ads, different OS a	pps, legacy
AWS Services	ECS, EKS, Fargate (for containers).		EC2, VMware	on AWS (for VMs)	
V	irtual machine	es		Containers	
VIRTUAL MACHINE	VIRTUAL MACHINE	VIRTUAL MACHINE	CONTAINER	CONTAINER	CONTAINER
АррА	Арр В	Арр С	Арр А	Арр В	Арр С
Bins/Libs	Bins/Libs	Bins/Libs	Bins/Libs	Bins/Libs	Bins/Libs
Guest OS	Guest OS	Guest OS		Container Engine	
	Hypervisor			Host Operating System	
	Infrastructure			Infrastructure	

13. Serverless compute plays a vital role in cloud computing environment. Justify with

reference to AWS Lamda Function.

Containors 1

Justification: Role of Serverless Compute in Cloud Computing (AWS Lambda)

AWS Lambda is a **serverless compute service** that allows running code without managing servers. It automatically scales, runs only when needed, and helps reduce costs.

Why Serverless Compute is Important?

✓ No Server Management

- No need to provision, maintain, or scale servers manually.
- AWS manages the infrastructure.

Auto Scaling

Fosturo

- Instantly scales up or down based on demand.
- Ideal for unpredictable workloads.

✓ Cost Efficiency

- Pay only for execution time (milliseconds).
- No cost when idle.

✓ Event-Driven Execution

- Runs on events (e.g., API requests, file uploads, database updates).
- Works with AWS services like S3, DynamoDB, API Gateway.

- Deploy functions quickly without worrying about infrastructure.
- Supports multiple languages (Node.js, Python, Java, etc.).

Example: AWS Lambda in Action

Imagine an image upload service:

- 1. A user uploads an image to **Amazon S3**.
- 2. An event triggers AWS Lambda.
- 3. Lambda resizes the image and stores it in another S3 bucket.

Benefit: No need for dedicated servers; the function runs **only when an image is uploaded**—saving costs and improving efficiency.

AWS Lambda **revolutionizes cloud computing** by eliminating server management, reducing costs, and enabling event-driven applications. It is a **key enabler of modern cloud architectures**, making applications more scalable and efficient.

14. Discuss the importance of Block Storage and Object Storage in Cloud Computing

Architecture. How Elastic Block Storage can be attached to specific EC2 instance

running Linux/Windows Operating System?

Importance of Block Storage and Object Storage in Cloud Computing

Block Storage:

- Structure: Stores data in small, fixed-size blocks.
- **Use:** Ideal for high-performance tasks like running databases or operating systems.
- **Example:** AWS Elastic Block Store (EBS) provides storage for EC2 instances.

Object Storage:

- **Structure:** Stores data as objects (data + metadata + unique ID).
- **Use:** Best for large files like photos, videos, and backups.
- **Example:** AWS S3 stores and manages large amounts of unstructured data.

How to Attach Elastic Block Storage (EBS) to EC2 Instances

For Linux EC2:

1. Create an EBS Volume:

• Go to AWS EC2 Dashboard \rightarrow Volumes \rightarrow Create Volume \rightarrow Set size and type.

2. Attach the EBS Volume to EC2 Instance:

Select the volume → Actions → Attach Volume → Select EC2 instance.

3. SSH into EC2 and Format the Volume:

 Connect using SSH → List devices with lsblk → Format the volume with sudo mkfs -t ext4 /dev/xvdf.

4. Mount the Volume:

Create a mount folder (sudo mkdir /mnt/data) → Mount volume (sudo mount /dev/xvdf /mnt/data).

For Windows EC2:

1. Create and Attach the EBS Volume:

• Follow similar steps to create and attach the volume to EC2.

2. Connect using RDP and Initialize the Volume:

• Open Disk Management → Initialize the disk.

3. Create and Format the Volume:

Create a new volume and assign a letter (e.g., D:) → Format the volume as NTFS.

4. Access the Volume:

• The volume appears in "This PC" with the assigned letter.

By using **Elastic Block Storage** with EC2, you can ensure your data is stored securely and can be accessed efficiently for your applications.

15. Your company have two VPC running in different AWS Regions. Each VPC has

Private and Public Subnets. You are required to exchange information between

instances running in Private Subnets. Design the VPC Peering with appropriate $% \left(1\right) =\left(1\right) +\left(1$

services.

Sure! Here's a simplified and easy-to-remember version of the VPC Peering setup process:

VPC Peering Overview

VPC Peering lets two VPCs (in same or different regions) securely communicate with each other, allowing instances in their private subnets to exchange information.

Steps to Set Up VPC Peering:

Step 1: Create Two VPCs

• Create VPC A and VPC B in AWS (public and private subnets in each).

Step 2: Create Peering Connection

- 1. In **VPC A**, go to **Peering Connections** and click **Create Peering Connection**.
- 2. Choose VPC A as Requestor and VPC B as Acceptor.
- 3. Create Connection and then go to VPC B to Accept it.

Step 3: Update Route Tables

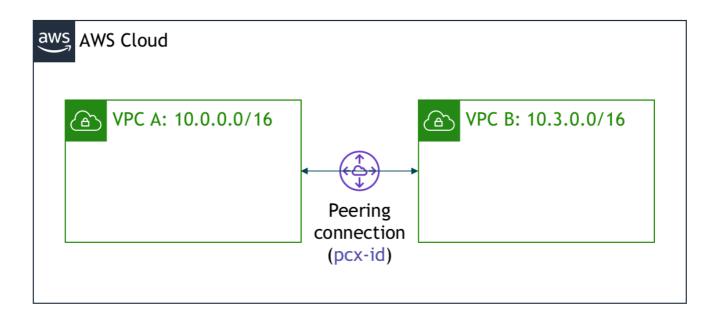
- 1. In **VPC A**, add a route to **VPC B's CIDR block**.
- 2. In **VPC B**, add a route to **VPC A's CIDR block**.

Step 4: Configure Security & ACLs

- 1. Update **Security Groups** to allow traffic between VPCs.
- 2. Ensure **Network ACLs** allow communication between the two VPCs.

Step 5: Test the Connection

• Launch EC2 instances in both VPCs and test if they can communicate (e.g., ping each other).



Route Table for VPC A

Destination	Target
10.0.0.0/16	local
10.3.0.0/16	pcx-id

Route Table for VPC B

Destination	Target
10.3.0.0/16	local
10.0.0.0/16	pcx-id

aws

Key Takeaways:

- Create VPCs: Set up two VPCs.
- Create Peering: Connect them.
- Update Routes: Allow traffic between VPCs.
- Security & ACLs: Ensure proper access.
- **Test**: Verify the connection is working.

16. Your company has on-premises Data Center and Cloud Data Recovery Center. The

private network of on-premises Data Center need to communicate with private

subnet of Cloud Data Recovery Center. Design the architecture of Site to Site $\ensuremath{\mathsf{VPN}}$

in AWS Cloud and on-premises Data Center to enable secured private communication between two private subnets.

To establish a **Site-to-Site VPN** between your **on-premises Data Center** and the **Cloud Data Recovery Center** in AWS, follow this simplified design:

Architecture Overview:

1. On-Premises Data Center:

• **Customer Gateway (CGW):** Your on-premises VPN device (e.g., router/firewall) to connect to AWS.

2. AWS Cloud Data Recovery Center:

- Virtual Private Cloud (VPC): Where your private subnet in AWS resides.
- **Virtual Private Gateway (VGW):** AWS-managed VPN device on the cloud side to enable secure connection.
- **VPN Connection:** The encrypted tunnel between the **CGW** and **VGW** for secure communication.

Steps to Set Up the Site-to-Site VPN:

1. Create Virtual Private Gateway (VGW) in AWS:

- Go to the AWS VPC console and create a Virtual Private Gateway.
- Attach this VGW to the VPC containing your private subnet.

2. Create Customer Gateway (CGW) for On-Premises:

- Create a Customer Gateway in AWS and enter the public IP address of your on-premises VPN device.
- Choose between **static routing** or **dynamic routing (BGP)** for the VPN connection.

3. Set Up VPN Connection:

- Create a **VPN connection** between the **VGW** and **CGW**.
- Download the configuration file from AWS (contains IPsec settings for your on-premises device).

4. Configure On-Premises VPN Device:

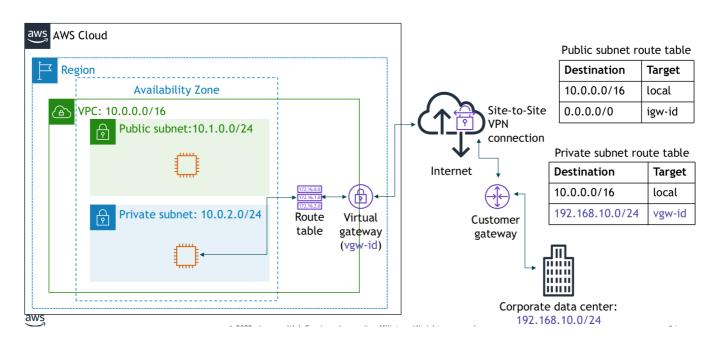
• Use the downloaded configuration to set up the VPN tunnel on your on-premises device.

5. Update Route Tables:

• Modify the route tables for both the AWS VPC (private subnet) and your on-premises network to route traffic through the VPN tunnel.

6. Test the VPN Connection:

• Verify that instances in the AWS private subnet can communicate with on-premises resources and vice versa.



Summary of the Steps:

- 1. Create VGW in AWS
- 2. Create CGW in AWS
- 3. Set up VPN connection
- 4. Configure on-premises VPN device
- 5. Update route tables
- 6. Test the connection

17. Discuss the role of Security Group in a Cloud Computing Environment with a Suitable Example.

Here's a simplified and easy-to-understand explanation of Security Groups in AWS, broken down into key points:

Role of Security Group in Cloud Computing:

1. Definition:

- A **Security Group** acts as a virtual firewall for your EC2 instances in the cloud.
- It **controls inbound** (incoming) and **outbound** (outgoing) traffic to/from the instance.
- Security Groups are **stateful**, meaning if you allow inbound traffic, the response traffic is automatically allowed.

2. How it Works:

- Security Groups allow you to set **rules** that either **allow** or **deny** access based on the **IP address** and **port** number.
- Inbound traffic needs to be explicitly allowed; if no rule exists, the traffic is denied by default.

3. Example:

- A web application hosted on an EC2 instance needs to:
 - Allow **HTTPS** traffic from the internet.
 - Allow **SSH** access from a specific IP range.

Steps to Set Up Security Group for EC2:

1. Create Security Group:

- Go to EC2 Dashboard in the AWS Management Console.
- In the **left sidebar**, click on **Security Groups** under the **Network & Security** section.
- Click on **Create Security Group**.

2. Configure Inbound Rules:

• Allow HTTPS Traffic:

Type: HTTPSProtocol: TCPPort Range: 443

• Source: **0.0.0.0/0** (this allows access from anywhere)

• Allow SSH Access from Specific IP:

Type: SSHProtocol: TCPPort Range: 22

• Source: **[Your IP Range]** (e.g., 192.168.1.0/24 for a specific range)

3. Configure Outbound Rules (if needed):

• Allow traffic to other destinations as required, for example:

Type: All trafficProtocol: All trafficPort Range: All

Destination: 0.0.0.0/0 (for unrestricted outbound access).

4. Attach Security Group to EC2 Instance:

- When launching your EC2 instance, **attach** the security group to the instance during the setup process.
- Alternatively, if the instance is already launched, **modify** its security group settings under the "Network & Security" section of the EC2 dashboard.

Summary of Security Group Example:

- The **web application** hosted on EC2 allows **HTTPS traffic** from anywhere but **restricts SSH** access to a specific IP range for security.
- Security groups **act as a firewall**, controlling who can access the application and how they can access it.

 ${f 18}.$ Discuss the role of Load Balancer in Cloud Computing Environment. Design the

Load Balancer with AWS ELB with a suitable example.

Role of Load Balancer in Cloud Computing:

- **Distributes Traffic**: Balances incoming traffic across multiple servers to avoid overload on any single server.
- Improves Availability: Routes traffic to healthy servers, ensuring the application is always available.
- Scalability: Allows adding more servers as traffic increases, handling more users.
- Fault Tolerance: If one server fails, traffic is redirected to other healthy servers.
- Performance Optimization: Ensures efficient performance by only sending traffic to healthy servers.

Designing Load Balancer with AWS ELB:

1. Set Up EC2 Instances:

• Launch multiple EC2 instances where your app runs.

2. Create Load Balancer:

- Go to AWS Console > EC2 > Load Balancers > Create Load Balancer.
- Choose **Application Load Balancer** (for web apps).
- Select HTTP/HTTPS listeners.

3. Configure Target Group:

- Create a target group with your EC2 instances.
- Set health checks (e.g., /health).

4. Add EC2 Instances to Target Group:

• Register your EC2 instances to the target group.

5. Configure Listeners & Routing:

• Set up listener rules (e.g., forward HTTP traffic to the target group).

6. Test the Load Balancer:

• Use the DNS name provided by AWS to test (e.g., myapp-alb-xxxx.elb.amazonaws.com).

Example:

- 3 EC2 Instances in different availability zones.
- AWS ELB distributes traffic across them.
- If one instance goes down, traffic is automatically sent to healthy ones.

19. When you run your applications on Cloud, you want to ensure that your architecture can scale to handle changes in demand. Discuss how to automatically

scale your EC2 instances with Amazon EC2 Auto Scaling

Automatically Scale EC2 Instances with Amazon EC2 Auto Scaling:

- Scale Automatically: Adds/removes EC2 instances based on demand.
- **Keeps App Available**: Ensures app stays up by adjusting instance numbers when needed.
- Handles Changes in Demand: Scales up during high traffic and down during low traffic.
- Replaces Unhealthy Instances: Detects and replaces unhealthy instances automatically.

Scaling Options:

- Manual: You adjust the number of instances yourself.
- **Scheduled**: Scale at set times (e.g., peak hours).
- **Dynamic**: Scales automatically based on real-time demand (using CloudWatch).
- **Predictive**: Predicts future demand and scales ahead of time.

Key Components:

- Auto Scaling Group: A group of EC2 instances that scales together.
- Launch Configuration: Template for instance setup (e.g., instance type).
- Works with ELB: Distributes traffic to healthy instances.

In Short:

- Automatically adjust EC2 instances based on traffic.
- Maintain performance while keeping costs low.
- Scale in four ways: manual, scheduled, dynamic, predictive.

20. Discuss the importance of combining Load balancing with Auto Scaling in cloud

computing environment.

Importance of Combining Load Balancing with Auto Scaling in Cloud Computing:

- Enhanced Application Availability:
 - **Load Balancer** ensures traffic is evenly distributed across multiple servers.
 - **Auto Scaling** adds/removes instances based on demand, ensuring that there are always enough resources to handle traffic.
 - Together, they make sure your app remains available even during traffic spikes or server failures.
- Optimal Performance:

- **Load Balancer** routes traffic to the healthiest instances.
- **Auto Scaling** adjusts the number of instances based on real-time usage to handle more requests or scale down during low traffic.
- This combination ensures your app performs well under varying loads.

• Cost Efficiency:

- **Auto Scaling** ensures you're not over-provisioning resources, reducing unnecessary costs by scaling down during low demand.
- **Load Balancer** ensures traffic is only sent to healthy, running instances, preventing wasted resources.

• Fault Tolerance:

- If an EC2 instance becomes unhealthy, **Load Balancer** redirects traffic to healthy instances.
- **Auto Scaling** automatically replaces unhealthy instances, maintaining system stability and availability.

Automatic Handling of Traffic Surges:

- Auto Scaling adds more instances when there is a sudden surge in traffic.
- **Load Balancer** efficiently distributes the increased traffic across the newly added instances, preventing overloading any single instance.

Seamless Scaling:

• **Load Balancer** integrates with **Auto Scaling** to ensure that scaling is smooth, and new instances can seamlessly handle traffic without any service disruption.

In Short:

- **Load Balancing** ensures efficient traffic distribution.
- Auto Scaling adjusts the number of EC2 instances based on demand.
- Together, they provide better availability, optimized performance, and cost-efficiency.

21. How multiple containers can be deployed with the aid of Docker Compose in EC2

instance. Justify with a suitable example.

To deploy multiple containers on an EC2 instance using Docker Compose, you can follow these simple steps:

1. Launch an EC2 Instance:

- Create an EC2 instance (e.g., Ubuntu) through the AWS Management Console.
- Make sure the security group allows traffic on the necessary ports (e.g., 80 for HTTP, 443 for HTTPS).

2. Connect to Your EC2 Instance:

• SSH into the EC2 instance:

```
ssh -i your-key.pem ec2-user@your-ec2-public-dns
```

3. Install Docker and Docker Compose:

• Update the package list:

```
sudo apt-get update
```

• Install Docker:

```
sudo apt-get install -y docker.io
```

• Install Docker Compose:

```
sudo curl -L
"https://github.com/docker/compose/releases/download/1.29.2/docker-
compose-$(uname -s)-$(uname -m)" -o /usr/local/bin/docker-compose
sudo chmod +x /usr/local/bin/docker-compose
```

4. Create a Docker Compose File:

• Create a directory for your project and navigate to it:

```
mkdir myapp
cd myapp
```

• Create a file called docker-compose.yml with the following content:

```
version: "3"
services:
    web:
    image: node:14
    ports:
        - "80:3000"
    volumes:
        - ./app:/usr/src/app
    working_dir: /usr/src/app
    command: npm start
    depends_on:
        - db
    db:
    image: mongo:latest
```

```
volumes:
- db-data:/data/db
volumes:
db-data:
```

5. Deploy the Application:

• Run the following command to start the containers:

```
docker-compose up -d
```

• Check the status of the containers:

```
docker-compose ps
```

6. Access the Application:

• Open a browser and visit your EC2 public IP address to see the web application running.

Why This Works:

- **Simplifies Setup:** Docker Compose allows you to define all your containers in a single file, making it easier to manage.
- Consistency: Using Compose ensures that your containers are always set up the same way.
- **Scalability:** You can easily add more containers or scale services as needed.

22. Discuss Six Pillars of AWS Well-Architected Framework with a suitable example.

Got it! Here's a simplified version of the Six Pillars of the AWS Well-Architected Framework with three key points for each pillar:

1. Operational Excellence

- Monitor and improve: Continuously track system performance and make improvements.
- Automate processes: Use automation to reduce manual tasks and improve efficiency.
- Enhance operations: Ensure systems are designed to adapt and improve over time.

2. Security

- Control access: Use strong authentication and IAM roles to manage who can access what.
- **Encrypt data**: Protect data at rest and in transit with encryption.
- Monitor for threats: Continuously monitor systems for vulnerabilities and unauthorized access.

3. Reliability

- **Prepare for failure**: Design systems to recover quickly from failures.
- Automate backups: Ensure automatic backups and data recovery options.
- **Test recovery**: Regularly test your disaster recovery procedures.

4. Performance Efficiency

- Use right-sized resources: Choose the most suitable resources based on demand.
- Optimize for performance: Continuously evaluate and adjust for better performance.
- Scale as needed: Implement auto-scaling to adjust resources automatically with demand.

5. Cost Optimization

- Avoid overprovisioning: Use only the resources needed to avoid unnecessary costs.
- Monitor usage: Track and optimize your cloud usage to identify cost-saving opportunities.
- **Use pricing models**: Take advantage of Reserved Instances or Savings Plans to save on long-term costs.

6. Sustainability

- Optimize resource use: Efficiently use resources to reduce energy consumption.
- Use eco-friendly regions: Choose AWS regions that use renewable energy sources.
- Minimize waste: Delete unused resources and data to reduce environmental impact.

Example: Deploying a Web Application on AWS

Consider deploying a web application on AWS. Here's how each pillar applies:

- **Operational Excellence**: Implement monitoring and logging to gain insights into application performance and user behavior.
- **Security**: Use AWS Identity and Access Management (IAM) to control access to resources and encrypt sensitive data.
- **Reliability**: Deploy the application across multiple Availability Zones to ensure high availability and automatic failover.
- **Performance Efficiency**: Utilize Amazon CloudFront for content delivery to reduce latency and improve load times.
- **Cost Optimization**: Choose appropriate instance types and use Auto Scaling to adjust capacity based on demand, ensuring you only pay for what you use.
- **Sustainability**: Opt for serverless services like AWS Lambda to reduce the need for provisioning and managing servers, thereby lowering energy consumption.
- 23. Explain the importance of reliability metrics like MTTF, MTTR, and MTBF in the
 - context of cloud computing, supported by a relevant example.

1. Mean Time to Failure (MTTF):

- **Definition:** The average time a non-repairable system operates before it fails.
- Why it's important: Helps determine how long a system can run without failure, indicating reliability.
- In cloud computing: Used to assess hardware like disks or storage devices.
- **Example:** If a storage disk has an MTTF of 8 years, it means the disk is expected to last, on average, for 8 years before failing.

2. Mean Time Between Failures (MTBF):

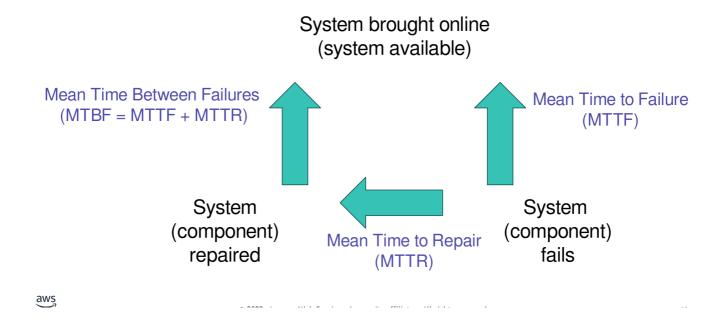
- **Definition:** The average time between failures in a repairable system.
- Why it's important: Indicates the frequency of failures, helping identify maintenance needs.
- In cloud computing: Used to monitor server uptime and cloud services to predict potential issues.
- **Example:** If a cloud server has an MTBF of 1,500 hours, it's expected to fail once every 1,500 hours of operation.

3. Mean Time to Repair (MTTR):

- **Definition:** The average time it takes to fix a system after a failure.
- Why it's important: Helps measure the efficiency of response and recovery from failures.
- In cloud computing: Helps improve service availability by tracking how quickly issues are resolved.
- **Example:** If a cloud service experiences an outage and it takes 30 minutes to fix, the MTTR is 30 minutes.

Importance in Cloud Computing:

- **Proactive Maintenance:** MTTF and MTBF help predict potential failures and guide preventive measures.
- Improved Service Availability: MTTR helps track how fast cloud services can recover, minimizing downtime.
- **Cost Efficiency:** By analyzing these metrics, cloud providers can optimize resource use and reduce unnecessary costs.
- **Enhanced User Experience:** A system with high MTTF/MTBF and low MTTR ensures reliable, smooth service for users.



24. Discuss the importance of availability tiers with a suitable example.

Sure! Here's a simplified version with easy-to-remember points:

Importance of Availability Tiers with Examples

1. Fault Tolerance

- What it means: Keeping things running even if something goes wrong.
- Why it's important: Availability tiers help keep your system working if one part fails by having backups in different locations.
- **Example:** In **Amazon RDS**, if the primary database fails, it switches to a backup in a different location, so the service keeps running smoothly.

2. Scalability

- What it means: The ability to grow and handle more users without changing the setup.
- Why it's important: Availability tiers let your system grow by adding more resources across different locations without needing a major redesign.
- Example: Elastic Load Balancing (ELB) helps spread traffic across different servers, so if more people use your app, it automatically adds more servers to handle the load.

3. Recoverability

- What it means: Quickly restoring the service if something breaks.
- Why it's important: Availability tiers make it easy to recover from big failures by using backup systems in different places.
- **Example:** In a web app, if one part of the system fails, the traffic is automatically sent to another working part, so users don't notice any issues.

Quick Recap:

- Fault Tolerance: Keep things running even if one part fails.
- Scalability: Easily add resources as your needs grow.
- **Recoverability:** Quickly restore services after a problem.

Availability	Max Disruption (per year)	Application Category
99%	3 days 15 hours	Batch processing, data extraction, transfer, and load jobs
99.9%	8 hours 45 minutes	Internal tools like knowledge management, project tracking
99.95%	4 hours 22 minutes	Online commerce, point of sale
99.99%	52 minutes	Video delivery, broadcast systems
99.999%	5 minutes	ATM transactions, telecommunications systems

aws

25. An organization aims to host multiple web services using containers. Explain how

Docker Compose can be utilized for deployment, including a relevant example.

To deploy multiple web services using containers in AWS with Docker Compose, follow these simple steps:

How Docker Compose Can Be Used for Deployment on AWS:

1. Docker Compose Overview:

- Docker Compose is a tool that allows you to define and manage multi-container Docker applications.
- It uses a YAML file (docker-compose.yml) to define all your services (containers), networks, and volumes needed for your application.

2. Set Up AWS Infrastructure:

- **EC2 Instance**: Start by creating an EC2 instance (virtual server) on AWS where Docker will run.
- Install Docker: SSH into the EC2 instance and install Docker and Docker Compose.

3. Create a docker-compose.yml File:

- In the docker-compose.yml file, define all the web services, such as front-end, back-end, and database services.
- Example:

```
version: "3"
services:
  web:
    image: nginx
  ports:
    - "80:80"
  db:
    image: mysql
    environment:
        MYSQL_ROOT_PASSWORD: password
```

4. Deploy Services Using Docker Compose:

• After defining the services, run the following command on your EC2 instance:

```
docker-compose up -d
```

• This will start all the containers (frontend, backend, and database) in detached mode.

5. Accessing Services:

- Ensure that your EC2 instance's security group allows incoming traffic on the ports your services are using (e.g., 80 for the front-end).
- You can access the web services by using the public IP of your EC2 instance.

Example:

Imagine the organization wants to host a web application with a front-end (React), back-end (Node.js), and a PostgreSQL database. The docker-compose.yml file can look like this:

```
version: "3"
services:
  web:
    image: nginx
  ports:
    - "80:80"
  db:
    image: mysql
  environment:
    MYSQL_ROOT_PASSWORD: password
```

Benefits of Using Docker Compose on AWS:

• **Simplifies Deployment**: With Docker Compose, you can define all your web services and dependencies in one file.

• **Ease of Scaling**: Docker Compose allows you to scale services easily by increasing the number of containers.

• **Consistency**: Using the same Docker Compose configuration in development and production ensures consistency across environments.